



Clever Girl: A Guide to Utilizing Color Histograms for Computer Vision and Image Search Engines

by **Adrian Rosebrock** on January 22, 2014 in **Image Descriptors**

Like 13

G+1 3

3



It's...it's a histogram. – Dr. Grant

Okay. So maybe that isn't the *exact* quote. But I think Dr. Grant would have been equally thrilled had he known the power of color histograms.

And more importantly, when the power goes out, *the histograms don't eat the tourists*.

So, what exactly is a histogram? A histogram represents the distribution of colors in an image. It can be visualized as a graph (or plot) that gives a high-level intuition of the intensity (pixel value) distribution. We are going to assume a RGB color space in this example, so these pixel values will be in the range of 0 to 255. If you are working in a different color space, the pixel range may be different.

When plotting the histogram, the X-axis serves as bins, then we are effectively counting the number use only 2 (equally spaced) bins, then we are col

Free 21-day crash course on computer vision & image search engines

[0, 128) or [128, 255]. The number of pixels binned to the x-axis value is then plotted on the y-axis.

Looking for the source code to this post?

[Jump right to the downloads section.](#)

OpenCV and Python versions:

This example will run on **Python 2.7** and **OpenCV 2.4.X/OpenCV 3.0+**.

By simply examining the histogram of an image, you get a general understanding regarding the contrast, brightness, and intensity distribution.

This post will give you an OpenCV histogram example.

Application to Image Search

In context of image search engines, histograms can be used to quantify an image and compare it to other images. In image search engines, we make the assumption that images are similar. I will talk more about this assumption in the future. For the time being, let's go ahead and assume the images are similar.

Comparing the "similarity" of color histograms can include: Euclidean, correlation, Chi-squared, etc. To use the Chi-squared distance, but the choice is up to you. No matter which distance metric you use, you can analyze histograms.

Free 21-day crash course on computer vision & image search engines

Interested in computer vision and image search engines, but don't know where to start? Let me help. I've created a free, 21-day crash course that is hand-tailored to give you the best possible introduction to computer vision. Sound good? Enter your email below to start your journey to becoming a computer vision master.

LET'S DO IT!

Touring Jurassic Park

Let's imagine that we were along with Dr. Grant and company on their first Jurassic Park tour. We brought along our cell phone to document the entire experience (and let's also pretend that camera phones were a "thing" back then). Assuming we didn't pull a Dennis Nedry and have our face eaten by a Dilophosaurus, we could later download the pictures from our smartphones to our computers and compute histograms for each of the images.

At the very beginning of the tour we spent a lot of time in the labs, learning about DNA and witnessing the hatching of a baby velociraptor. These labs have a lot of "steel" and "gray" colors to them. Later on, we got into our jeeps and drove into the park. The park itself is a jungle — lots of green colors.

So based on these two color distributions, which one do you think the Dr. Grant image above is more similar to?

Free 21-day crash course on computer vision & image search engines

Well, we see that there is a fair amount of greenery in the background of the photo. In all likelihood, the color distribution of the Dr. Grant photo would be more “similar” to our pictures taken during the jungle tour vs. our pictures taken in the lab.



You can master computer vision...*in a single weekend.*

[Click here to become a computer vision guru](#)

Using OpenCV to Compute Histograms

Now, let's start building some color histograms of

We will be using the `cv2.calcHist` function in `cv2`. In any code examples, let's quickly review the function

`cv2.calcHist(images, channels, mask,`

1. **images:** This is the image that we want to compute a histogram for. It is a `numpy.ndarray` of type `uint8` or `uint16` and can be a single image or a list of images (`[myImage]`).
2. **channels:** A list of indexes, where we specify the channels to compute a histogram for. To compute a histogram of a single channel, we pass the index of the channel. To compute a histogram for all three red, green, and blue channels, we pass `[0, 1, 2]`.
3. **mask:** I haven't covered masking yet in this tutorial, but it allows us to only compute a histogram for a portion of the image. The mask should have the same shape as our original image, with pixels with a value greater than zero indicating the area to include. If we pass `None`, it allows us to only compute a histogram for a portion of the image.
4. **histSize:** This is the number of bins we want to use when computing a histogram. Again, this is a list, one for each channel we are computing a histogram for. The bin sizes do not all have to be the same. Here is an example of 32 bins for each channel: `[32, 32, 32]`.
5. **ranges:** The range of possible pixel values. Normally, this is `[0, 256]` for each channel, but if you are using a color space other than RGB (such as HSV), the ranges might be different.

Now that we have an understanding of the `cv2.calcHist` function, let's write some actual code.

Load Our Image

Python

```
1 # import the necessary packages
2 from matplotlib import pyplot as plt
3 import numpy as np
4 import argparse
5 import cv2
6
7 # construct the argument parser and parse the arguments
8 ap = argparse.ArgumentParser()
9 ap.add_argument("-i", "--image", required = True, help = "Path to the image")
10 args = vars(ap.parse_args())
11
12 # load the image and show it
13 image = cv2.imread(args["image"])
14 cv2.imshow("image", image)
```

Free 21-day crash course on computer vision & image search engines

This code isn't very exciting yet. All we are doing is importing the packages we will need, setting up an argument parser, and loading our image.

Computing a Grayscale Histogram using Python and OpenCV	Python
<pre> 1 # convert the image to grayscale and create a histogram 2 gray = cv2.cvtColor(image, cv2.COLOR_BGR2GRAY) 3 cv2.imshow("gray", gray) 4 hist = cv2.calcHist([gray], [0], None, [256], [0, 256]) 5 plt.figure() 6 plt.title("Grayscale Histogram") 7 plt.xlabel("Bins") 8 plt.ylabel("# of Pixels") 9 plt.plot(hist) 10 plt.xlim([0, 256]) </pre>	

Now things are getting a little more interesting. On line 2, we convert the image from the BGR colorspace to grayscale. Line 4 computes the actual histogram of the code up with the function documentation at [cv2.calcHist\(\)](#). A grayscale image has only one channel. We don't have a mask, so we set the mask to None, and the possible values range from 0 to 255.

A call to `plt.show()` displays:



Figure 1: Dr. Grant grayscale histogram.

Not bad. How do we interpret this histogram? Well, the bins (0-255) are plotted on the X-axis. And the Y-axis counts the number of pixels in each bin. The majority of pixels fall in the range of ~50 to ~125. Looking at the right tail of the histogram, we see very few pixels in the range 200 to 255. This means that there are very few "white" pixels in the image.

Now that we've seen at a grayscale histogram, let's look at what I call a "flattened" color histogram:

Computing a Flattened Color Histogram using Python and OpenCV	Python
<pre> 1 # grab the image channels, initialize the tuple of colors, 2 # the figure and the flattened feature vector 3 chans = cv2.split(image) 4 colors = ("b", "g", "r") 5 plt.figure() 6 plt.title("'Flattened' Color Histogram") 7 plt.xlabel("Bins") 8 plt.ylabel("# of Pixels") 9 features = [] 10 11 # loop over the image channels 12 for (chan, color) in zip(chans, colors): 13 # create a histogram for the current </pre>	

Free 21-day crash course on computer vision & image search engines

Interested in computer vision and image search engines, but don't know where to start? Let me help. I've created a free, 21-day crash course that is hand-tailored to give you the best possible introduction to computer vision. Sound good? Enter your email below to start your journey to becoming a computer vision master.

LET'S DO IT!

Free 21-day crash course on computer vision & image search engines

```

14 # concatenate the resulting histograms for each
15 # channel
16 hist = cv2.calcHist([chan], [0], None, [256], [0, 256])
17 features.extend(hist)
18
19 # plot the histogram
20 plt.plot(hist, color = color)
21 plt.xlim([0, 256])
22
23 # here we are simply showing the dimensionality of the
24 # flattened color histogram 256 bins for each channel
25 # x 3 channels = 768 total values -- in practice, we would
26 # normally not use 256 bins for each channel, a choice
27 # between 32-96 bins are normally used, but this tends
28 # to be application dependent
29 print "flattened feature vector size: %d" % (np.array(features).flatten().shape)

```

There's definitely more code involved in computing a color histogram. Let's tear this code apart and get a better understanding of what's going on.

- **Lines 3-4:** The first thing we are going to do is to read the image in BGR (Blue, Green, and Red). Normally, we read this is a NumPy array in reverse order: BGR. The tuple of strings representing the colors.
- **Lines 5-9:** Here we are just setting up our PyPlot figure and axes to display the histograms.
- **Line 12:** Let's start looping over the channels.
- **Line 16-17:** We are now computing a histogram for each channel. As computing a histogram for a single channel, we append the result to our features list.
- **Line 20:** Plot the histogram using the current color.
- **Line 29:** Here we are just examining the shape of the "flattened" histogram not because the (1) histogram is flattened using the flatten() method. I call this a "flattened" histogram because it's a 1D array of counts. Later, we explore multi-dimensional histograms (2D and 3D). A flattened histogram is simply the histogram *for each individual channel concatenated together*.


Now let's plot our color histograms:

Free 21-day crash course on computer vision & image search engines

Interested in computer vision and image search engines, but don't know where to start? Let me help. I've created a free, 21-day crash course that is hand-tailored to give you the best possible introduction to computer vision. Sound good? Enter your email below to start your journey to becoming a computer vision master.

LET'S DO IT!

Free 21-day crash course on computer vision & image search engines



Awesome. That was pretty simple. What does this range of blue pixel values around bin #50. This range of blue range of green pixels from bin #50 to #125 refers

Multi-dimensional Histograms

Up until this point, we have computed a histogram for each color channel. Now we will move on to multi-dimensional histograms and take into consideration two channels at a time.

The way I like to explain multi-dimensional histograms is to use the word **AND**. For example, we can ask a question such as

“how many pixels have a Red value of 10 **AND** a Blue value of 30?” How many pixels have a Green value of 200 **AND** a Red value of 130? By using the conjunctive **AND** we are able to construct multi-dimensional histograms.

It's that simple. Let's checkout some code to automate the process of building a 2D histogram:

Computing a 2D Color Histogram using Python and OpenCV	Python
<pre> 1 # let's move on to 2D histograms -- I am reducing the 2 # number of bins in the histogram from 256 to 32 so we 3 # can better visualize the results 4 fig = plt.figure() 5 6 # plot a 2D color histogram for green and blue 7 ax = fig.add_subplot(131) 8 hist = cv2.calcHist([chans[1], chans[0]], [0, 1], None, 9 [32, 32], [0, 256, 0, 256]) 10 p = ax.imshow(hist, interpolation = 'nearest') 11 ax.set_title("2D Color Histogram for Green and Blue") 12 plt.colorbar(p) </pre>	<p>Free 21-day crash course on computer vision & image search engines</p> <p>Interested in computer vision and image search engines, but don't know where to start? Let me help. I've created a free, 21-day crash course that is hand-tailored to give you the best possible introduction to computer vision. Sound good? Enter your email below to start your journey to becoming a computer vision master.</p> <p>Email Address</p> <p>LET'S DO IT!</p>

```

13
14 # plot a 2D color histogram for green and red
15 ax = fig.add_subplot(132)
16 hist = cv2.calcHist([chans[1], chans[2]], [0, 1], None,
17     [32, 32], [0, 256, 0, 256])
18 p = ax.imshow(hist, interpolation = "nearest")
19 ax.set_title("2D Color Histogram for Green and Red")
20 plt.colorbar(p)
21
22 # plot a 2D color histogram for blue and red
23 ax = fig.add_subplot(133)
24 hist = cv2.calcHist([chans[0], chans[2]], [0, 1], None,
25     [32, 32], [0, 256, 0, 256])
26 p = ax.imshow(hist, interpolation = "nearest")
27 ax.set_title("2D Color Histogram for Blue and Red")
28 plt.colorbar(p)
29
30 # finally, let's examine the dimensionality of
31 # the 2D histograms
32 print "2D histogram shape: %s, with %d values" %
33     hist.shape, hist.flatten().shape[0])

```

Yes, this is a fair amount of code. But that's only for each combination of RGB channels: Red and Green.

Now that we are working with multi-dimensional histograms, the bins we are using. In previous examples, I've used 256 bins for each dimension in a 2D histogram. We used a 256 bins for each dimension in a 2D histogram. We have 65,536 separate pixel counts. Not only is this a lot of data, but applications using somewhere between 8 and 64 bins per dimension. As **Lines 8-9** show, I am now using 32 bins for each dimension.

The most important take away from this code can be the `cv2.calcHist` function. Here we see that we are using the Red and Blue channels. And that's all there is to it.

So how is a 2D histogram stored in OpenCV? It's a 2D NumPy array. Since I used 32 bins for each channel, I now have a 32×32 histogram. We can treat this histogram as a feature vector simply by flattening it (**Line 33**). Flattening our histograms yields a list with 1024 values.

How do we visualize a 2D histogram? Let's take a look.

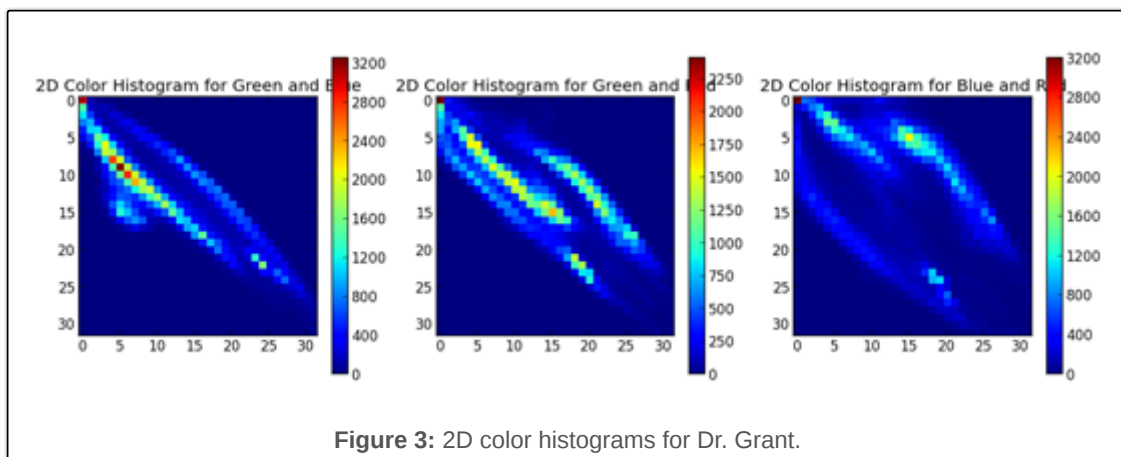


Figure 3: 2D color histograms for Dr. Grant.

In the above Figure, we see three graphs. The first graph shows the 2D color histogram for the Red and Green channels, the second for Green and Red, and the third for Blue and Red.

Free 21-day crash course on computer vision & image search engines

Interested in computer vision and image search engines, but don't know where to start? Let me help. I've created a free, 21-day crash course that is hand-tailored to give you the best possible introduction to computer vision. Sound good? Enter your email below to start your journey to becoming a computer vision master.

LET'S DO IT!

Free 21-day crash course on computer vision & image search engines

low pixel counts, whereas shades of red represent large pixel counts (i.e. peaks in the 2D histogram). We can see such a peak in the Green and Blue 2D histogram (the first graph) when $X=5$ and $Y=10$.

Using a 2D histogram takes into account two channels at a time. But what if we wanted to account for *all three* RGB channels? You guessed it. We're now going to build a 3D histogram.

Computing a 3D Color Histogram using Python and OpenCV	Python
<pre> 1 # our 2D histogram could only take into account 2 out 2 # of the 3 channels in the image so now let's build a 3 # 3D color histogram (utilizing all channels) with 8 bins 4 # in each direction -- we can't plot the 3D histogram, but 5 # the theory is exactly like that of a 2D histogram, so 6 # we'll just show the shape of the histogram 7 hist = cv2.calcHist([image], [0, 1, 2], 8 None, [8, 8, 8], [0, 256, 0, 256, 0, 9 print "3D histogram shape: %s, with %d va 10 hist.shape, hist.flatten().shape[0]) </pre>	

The code here is very simple — it's just an extension of an 8x8x8 histogram for each of the RGB channels. We can see that the shape is indeed (8, 8, 8) with 512 values. This can be done by simply flattening the array.

Color Spaces

The examples in this post have only explored the RGB color space constructed for any color space in OpenCV. Discuss this in the next post, but if you are interested, check out the [documentation](#).

Drawbacks

Earlier in this post we made the assumption that color histograms are semantically similar. For small, simple datasets, this may in fact be true. However, in practice, this assumption does not always hold.

Let's think about why this for.

For one, color histograms, by definition ignore both the shape and texture of the object(s) in the image. This means that color histograms have no concept of the shape of an object or the texture of the object. Furthermore, histograms also disregard any spatial information (i.e. where in the image the pixel value came from). An extension to the histogram, the color correlogram, can be used to encode a spatial relationship amongst pixels.

Let's think about [Chic Engine](#), my visual fashion search engine iPhone app. I have different categories for different types of clothes, such as shoes and shirts. If I were using color histograms to describe a red shoe and a red shirt, the histogram would assume they were the same object. Clearly they are both red, but the semantics end there — they are simply not the same. Color histograms simply have no way to “model” what a shoe or a shirt is.

Finally, color histograms are sensitive to “noise”, and the image was captured under and quantization error.

Free 21-day crash course on computer vision & image search engines

Interested in computer vision and image search engines, but don't know where to start? Let me help. I've created a free, 21-day crash course that is hand-tailored to give you the best possible introduction to computer vision. Sound good? Enter your email below to start your journey to becoming a computer vision master.

LET'S DO IT!

Free 21-day crash course on computer vision & image search engines

limitations can potentially be mitigated by using a different color space than RGB (such as HSV or L*a*b*).

However, all that said, histograms are still widely used as image descriptors. They are dead simple to implement and very fast to compute. And while they have their limitations, they are very powerful when used correctly and in the right context.

What's next?

In this post I provided an OpenCV histogram example using Python, from start to finish.

Coming up Monday, we are going to build our first histograms. Are you excited? I am. And maybe I've been building up around here. No guarantees tho.

Until Monday!

Downloads:



If you would like to download your email address in the form also send you a **FREE 11-page** Image Search Engines, including **exclusive** techniques. Sound good? If so, enter your email address immediately!

Email address:

DOWNLOAD THE CODE!

Free 21-day crash course on computer vision & image search engines

Interested in computer vision and image search engines, but don't know where to start? Let me help. I've created a free, 21-day crash course that is hand-tailored to give you the best possible introduction to computer vision. Sound good? Enter your email below to start your journey to becoming a computer vision master.

LET'S DO IT!

Resource Guide (it's totally free).



Enter your email address below to get my **free 11-page Image Search Engine Resource Guide PDF**. Uncover **exclusive techniques** that I don't publish on this blog and start building image search engines of your own!

DOWNLOAD THE GUIDE!

Free 21-day crash course on computer vision & image search engines

2d histogram, 3d histogram, color, histogram, rgb

< Basic Image Manipulations in Python and OpenCV : Resizing (scaling), Rotating, and Cropping

Hobbits and Histograms – A How-To Guide to Building Your First Image Search Engine in Python >

47 Responses to *Clever Girl: A Guide to Utilizing Color Histograms for Computer Vision and Image Search Engines*



mamachanko February 1, 2014 at 3:59 am #

I think you meant to say that the common range is [0, 256] when describing the calcHist function.



mamachanko February 1, 2014 at 4:01 am #

Indeed, it seems I was wrong.



Adrian Rosebrock February 1, 2014 at 4:01 am #

This used to throw me for a loop when I was learning OpenCV and NumPy “gotchas” I’ve encountered.

So yes, the possible pixel values for each channel are 0-255. However, when using calcHist (and NumPy’s histogram method) the range becomes [0, 255). (Notice the parenthesis, not a bracket). It’s an extremely subtle difference, and in most cases, it wouldn’t “break” anything. But in order to capture the range [0, 255] with 256 bins, we have to specify [0, 256). It’s definitely annoying.

Free 21-day crash course on computer vision & image search engines

Interested in computer vision and image search engines, but don't know where to start? Let me help. I've created a free, 21-day crash course that is hand-tailored to give you the best possible introduction to computer vision. Sound good? Enter your email below to start your journey to becoming a computer vision master.

LET'S DO IT!



Jeremy November 22, 2014 at 3:10 pm #

REPLY ↩

I had a question on line 9: `ap.add_argument("-i", "--image", required = True, help = "Path to the image")`

If my path is: `C:\Users\Jeremy\Documents\IPython Notebooks`, would my line 9 look like `ap.add_argument("IPython Notebooks", "face.jpeg", required = True, help = "Path to the image")`?



Adrian Rosebrock November 23, 2014 at 10:21 am #

Free 21-day crash course on computer vision & image search engines

Hi Jeremy. If you are using an IPython Notebook you can skip the argument parsing entirely and use the path to your image directly. The argument parsing is only necessary if you are executing the script via command line.



Phil GI October 31, 2015 at 2:46 am #

REPLY ↩

I'm attempting to plot a 3D histogram. I thought it would be simple extension of the 2D examples you provide using something like:

```
hist = cv2.calcHist([img_patch],[0,1,2],None,[8,8,8],[0,256,0,256,0,256])
```

Where the Z axis would provide 'red' information respectively.

Not having much luck with it. Any guidance would



Shelly November 14, 2015 at 3:48 pm #

Did you have any luck with this? See



Mohanish December 15, 2015 at 9:10 pm #

hey, iam getting this error

```
1 "usage: mycv_3.py [-h] -i IMAGE
2 mycv_3.py: error: argument -i/--image i
```

iam really stucked in it...

help me out !!!

Free 21-day crash course on computer vision & image search engines

Interested in computer vision and image search engines, but don't know where to start? Let me help. I've created a free, 21-day crash course that is hand-tailored to give you the best possible introduction to computer vision. Sound good? Enter your email below to start your journey to becoming a computer vision master.

LET'S DO IT!



Adrian Rosebrock December 16, 2015 at 6:37 am #

REPLY ↩

Hey Mohanish, you need to supply the `--image` switch, which is the path to the image you want to compute a histogram for:

```
$ python mycv_3.py --image path/to/my/image.jpg
```

I would suggest downloading the .zip of the source code to this post which includes the example image and example command to run.



Hamid December 16, 2015 at 4:54 am #

REPLY ↩

Dear Adrian,
Thank you for providing awesome material around

Free 21-day crash course on computer
vision & image search engines

histogram of images in a database, so I tried to append each of the feature list in another list as follow:

```
1 for (chan, color) in zip(chans, colors):
2     hist = cv2.calcHist([chan], [0], None, [256], [0, 256])
3     features.extend(hist)
4
5     Matrix.append(features)
```

but when I write the martix as a csv file, I get a [] around each number in the csv file, do you have any idea how to work out this problem, as I need to read color histograms as an excel file in matlab later.

I truly appreciate your kind guidance and help.

Thank you.



Adrian Rosebrock December 16, 2016

Hey Hamid, have you tried looking at my post demonstrates how to extract color histograms to a CSV file.



Stephen S Eacuellon June 23, 2016 at 11:04 am

Hi Adrian,

When I run the code with the image you provided

OpenCV Error: Assertion failed (size.width>0 && size.height>0) in function cv::cvtColor, file D:\Build\OpenCV\opencv-3.1.0\modules\highgui\src\cvtColor.cpp, line 236

Seems like openCV thinks the image size is 0? Strangely it runs with .png files correctly. Any idea what might cause an error like this?



Adrian Rosebrock June 23, 2016 at 1:04 pm #

REPLY ↩

It's a strange error, but it seems like that your system is not configured to load JPEG files. If it .png files are working correctly, then you have the proper PNG libraries installed. But I'm guessing that the JPEG libraries are not installed/are not working with OpenCV.



Farah July 13, 2016 at 6:14 am #

REPLY ↩

Hi Adrian

I wanted an equal density histogram for my research purpose. Is there some way of specifying the bins with calcHist rather than the size of the bin. I used 10 then the second may be from 11 to 16 and so on.

Free 21-day crash course on computer vision & image search engines



Adrian Rosebrock July 14, 2016 at 1:13 pm #

REPLY ↩

If you're looking for non-equal sized bins, I probably wouldn't use the `cv2.calcHist` function for this. Instead, the [NumPy's histogram function](#) will give you the fine-grained control that you're looking for.



Jooyeon July 21, 2016 at 12:18 pm #

REPLY ↩

Hello, Adrian!

I have downloaded the code and ran this code but I got an error. I have followed the tutorial from your blog post (<http://www.pyimagesearch.com/2015/06/15/installing-matplotlib-on-the-cv-virtual-environment>) on the cv virtual environment.

It has this runtime error: `RuntimeError: Python is not supported as a backend. Python backend will not be able to function correctly if PyOpenGL is installed. See http://www.pyimagesearch.com/2015/06/15/installing-matplotlib-on-the-cv-virtual-environment for more information on installing PyOpenGL. You can also try to reinstall Python as a framework, or try one of the other backends. For more information see 'Working with Matplotlib' in the documentation.`

I wonder how you solved this problem. Thank you!

Free 21-day crash course on computer vision & image search engines

Interested in computer vision and image search engines, but don't know where to start? Let me help. I've created a free, 21-day crash course that is hand-tailored to give you the best possible introduction to computer vision. Sound good? Enter your email below to start your journey to becoming a computer vision master.

LET'S DO IT!



Adrian Rosebrock July 21, 2016 at 12:18 pm #

Hey Jooyeon — this error happens incompatibly with the virtual environments. The developers are aware of the problem, but I don't know of an exact fix for the newer versions of matplotlib. Instead, I would recommend installing a previous version of matplotlib:

```
$ pip install matplotlib==1.4.3
```

This should take care of the issue.



Jooyeon July 22, 2016 at 2:37 am #

REPLY ↩

Thanks, Adrian. It works beautifully! Can't wait to dive into more openCV projects. I hope I can be an expert like you someday, sharing my knowledge with others to build awesome things!!



Adrian Rosebrock July 22, 2016 at 10:55 am #

REPLY ↩

Awesome, I'm glad to hear it!

Free 21-day crash course on computer vision & image search engines

**AMAN KHATRI** August 5, 2016 at 5:56 am #

REPLY ↩

I have installed 1.4.3 version but the above mentioned error is still there for me.what can i do to remove the error?

**Adrian Rosebrock** August 7, 2016 at 8:20 am #

REPLY ↩

You might want to double-check which version of matplotlib you are using and whether or not you are using a Python virtual environments are independent switched matplotlib versions for the ir

Free 21-day crash course on computer vision & image search engines

Interested in computer vision and image search engines, but don't know where to start? Let me help. I've created a free, 21-day crash course that is hand-tailored to give you the best possible introduction to computer vision. Sound good? Enter your email below to start your journey to becoming a computer vision master.

LET'S DO IT!

**Devi** September 19, 2016 at 3:59 am #

Thank you for sharing. Nicely explained

**Adrian Rosebrock** September 19, 2016 at 4:00 am #

Fantastic, I'm happy to hear the tuto

**chaidev** October 18, 2016 at 11:46 am #

Adrian, I just wanted to tell you how thoroughly impressed I am with this blog and with your commitment to keep answering comments on (older) posts. I'll be following along as i'm in looking to apply some CV in upcoming projects.

The threshold for following your tuts is low, yet the extra info along the way really helps build context for what we're actually doing.

thank you so much for your effort!

**Adrian Rosebrock** October 20, 2016 at 8:51 am #

REPLY ↩

Thank you for the kind words Chaidev, it's comments like these that keep me coming back to answer comments on older posts 😊

**Dipti** November 26, 2016 at 3:16 pm #

How to draw a histogram of an RGB ima

Free 21-day crash course on computer vision & image search engines



Adrian Rosebrock November 28, 2016 at 10:31 am #

REPLY ↩

I would use `matplotlib` to draw the histograms. I cover computing and extracting histograms in detail inside my book, [Practical Python and OpenCV](#).



Vishnu January 19, 2017 at 8:04 am #

REPLY ↩

Hi Adrian!

So Im realtively new to python and Im not able to
what is the purpose of lines 8-10? What does it do

**Free 21-day crash course
on computer vision &
image search engines**



Interested in computer vision and image search engines, but don't know where to start? Let me help. I've created a free, 21-day crash course that is hand-tailored to give you the best possible introduction to computer vision. Sound good? Enter your email below to start your journey to becoming a computer vision master.

Email Address

LET'S DO IT!



Adrian Rosebrock January 20, 2017 at 10:04 am #

Hey Vishnu — I would suggest you s
[arguments](#).



S Rank February 25, 2017 at 6:18 pm #

If anyone is getting the following error w
10.11.5):

Terminating app due to uncaught exception 'NSIn
_setup:]: unrecognized selector sent to instance C

This is what I did to resolve it:

```
pip install PyQt5
```

(at the very top of the code file)

```
import matplotlib  
matplotlib.use("Qt5Agg")
```

I understand the problem is matplotlib relies on the OS for certain image processing/rendering functionality, and on OSX it tries to interface with the OS in an unsupported way. By specifying exactly how to interface with the OS, this problem is resolved.

Ctrl + F tags: MacOS, Apple



Adrian Rosebrock February 27, 2017 at 11:14 am #

REPLY ↩

Thanks for sharing!

**Free 21-day crash course on computer
vision & image search engines**



REPLY 21



REPLY ↩



Free 21-day crash course on computer vision & image search engines



Email Address



REPLY ↩

REPLY 

REPLY ↩

Free 21-day crash course on computer vision & image search engines

which would represent an images histogram for example. I had tried cv2.calcHist but I'm struggling to pull it together into code. Any help appreciated!



Adrian Rosebrock May 21, 2017 at 5:18 am #

REPLY ↩

Hi Richie — you're in luck, I'll be publishing a blog post on how to determine the "most colorful" and "least colorful" images in a dataset later this month. Be sure to keep your eye on the PyImageSearch blog!

Trackbacks/Pingbacks

[Hobbits and Histograms - A How-To Guide to Building an Image Search Engine - PyImageSearch](#) - January 27, 2014

[...] discussed the color histogram in my previous post on computer vision and image search engines. If you haven't read it yet, please go back and read it after [...]

[URL](#) - January 29, 2014

... **[Trackback]**

[...] Read More here: pyimagesearch.com/2014/01/22/clever-girl-a-guide-to-utilizing-color-histograms-for-computer-vision-and-image-search-engines/ [...]

[Building an Image Search Engine: Indexing Your Data](#) - January 17, 2014

[...] Color: Image descriptors that characterize the color of an image. These include the mean, standard deviation, and skewness, along with color histograms.

[Building an Image Search Engine: Defining Your Similarity Metric](#) - January 17, 2014

[...] what it takes to build an image search engine. Start by extracting features from images using simple color histograms. Then compare them using the distance functions discussed above. Note your [...]

[How To Describe and Quantify an Image Using Feature Vectors](#) - March 3, 2014

[...] back to the Clever Girl: A Guide to Utilizing Color Histograms for Computer Vision and Image Search Engines and Hobbits and Histograms, we could also use a 3D color histogram to describe our [...]

[Building an Image Search Engine: Defining Your Image Descriptor \(Step 1 of 4\) - PyImageSearch](#) - April 28, 2014

[...] the mean, standard deviation, and skew of each channel's pixel intensities. We could also use color histograms as we've seen in other blog posts. In color histograms are global image descriptors applied [...]

[Color Quantization with OpenCV using K-Means Clustering](#) - PyImageSearch - July 7, 2014

[...] given 24-bit RGB image has 256 x 256 x 256 possible colors. And sure, we can build standard color histograms based on these intensity [...]

[How-To: 3 Ways to Compare Histograms using OpenCV and Python](#) - PyImageSearch - July 14, 2014

[...] For more details on the cv2.calcHist function, definitely take a look at my guide to utilizing color histograms for computer vision and image search engines post. [...]

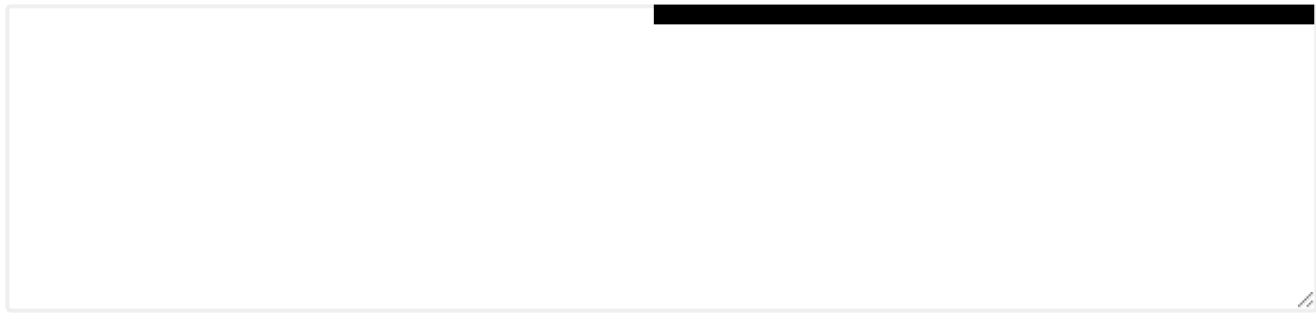
Free 21-day crash course on computer vision & image search engines

Interested in computer vision and image search engines, but don't know where to start? Let me help. I've created a free, 21-day crash course that is hand-tailored to give you the best possible introduction to computer vision. Sound good? Enter your email below to start your journey to becoming a computer vision master.

LET'S DO IT!

Leave a Reply

Free 21-day crash course on computer vision & image search engines



Name (required)

Email (will not be published) (required)

Website

SUBMIT COMMENT

Resource Guide (it's totally free).



Click the button below to get my **free** **Resource Guide**.
Uncover **exclusive techniques** that I've used to build my own image search engines of your own.

Download

Free 21-day crash course on computer vision & image search engines



Interested in computer vision and image search engines, but don't know where to start? Let me help. I've created a free, 21-day crash course that is hand-tailored to give you the best possible introduction to computer vision. Sound good? Enter your email below to start your journey to becoming a computer vision master.

LET'S DO IT!

Deep Learning for Computer Vision with Python Book




You're interested in deep learning and computer vision, then this **book will teach you all you need to know about deep learning**.

Free 21-day crash course on computer vision & image search engines

w

[CLICK HERE TO PRE-ORDER MY NEW BOOK](#)

You can detect faces in images & video.



Are you interested in **detecting faces in images & video**? Then let me help! I guarantee that my new book will turn [Click here to give it a shot yourself.](#)

[CLICK HERE TO MASTER FACE DETECTION](#)

PyImageSearch Gurus: NOW ENROLLING!

Free 21-day crash course on computer vision & image search engines

Interested in computer vision and image search engines, but don't know where to start? Let me help. I've created a free, 21-day crash course that is hand-tailored to give you the best possible introduction to computer vision. Sound good? Enter your email below to start your journey to becoming a computer vision master.

LET'S DO IT!

The PyImageSearch Gurus course is *now enrolling!* Inside the course you'll learn how to perform:

- Automatic License Plate Recognition (ANPR)
- Deep Learning
- Face Recognition
- *and much more!*

Click the button below to learn more about the course, take a tour, and get 10 (FREE) sample lessons.

[TAKE A TOUR & GET 10 \(FREE\) LESSONS](#)


Hello! I'm Adrian Rosebrock.

**Free 21-day crash course on computer
vision & image search engines**



I'm an entrepreneur and Ph.D who has launched two successful image search engines, [ID My Pill](#) and [Chic Engine](#). I'm here to share my tips, tricks, and hacks I've learned along the way.

Learn computer vision in a single weekend.



Want to learn computer vision & OpenCV? I can teach you in a single weekend. My new book is your **guaranteed, quick-start guide** to becoming a computer vision master. Click here to become a computer vision ninja.

[CLICK HERE TO BECOME AN OPENCV NINJA](#)

Free 21-day crash course on computer vision & image search engines

Interested in computer vision and image search engines, but don't know where to start? Let me help. I've created a free, 21-day crash course that is hand-tailored to give you the best possible introduction to computer vision. Sound good? Enter your email below to start your journey to becoming a computer vision master.

LET'S DO IT!

Subscribe via RSS



Never miss a post! Subscribe to the PyImageSearch RSS Feed and keep up to date with my image search engine tutorials, tips, and tricks

POPULAR

Install OpenCV and Python on your Raspberry Pi 2 and B+
FEBRUARY 23, 2015

Home surveillance and motion detection with the Raspberry Pi, Python, OpenCV, and Dropbox
JUNE 1, 2015

Install guide: Raspberry Pi 3 + Raspbian Jessie + OpenCV 3
APRIL 18, 2016

How to install OpenCV 3 on Raspbian Jessie
OCTOBER 26, 2015

Basic motion detection and tracking with Python and OpenCV
MAY 25, 2015

Free 21-day crash course on computer vision & image search engines

Accessing the Raspberry Pi Camera with OpenCV and Python

MARCH 30, 2015

Install OpenCV 3.0 and Python 2.7+ on Ubuntu

JUNE 22, 2015

Search



Find me on [Twitter](#), [Facebook](#), [Google+](#), and [LinkedIn](#).
© 2017 PyImageSearch. All Rights Reserved.

Free 21-day crash course on computer vision & image search engines



Interested in computer vision and image search engines, but don't know where to start? Let me help. I've created a free, 21-day crash course that is hand-tailored to give you the best possible introduction to computer vision. Sound good? Enter your email below to start your journey to becoming a computer vision master.

LET'S DO IT!

Free 21-day crash course on computer
vision & image search engines